

Pydantic x Modal / October 2026

Why even have an SDK?

Samuel Colvin · Pydantic

whoami - Samuel Colvin

- Creator of **Pydantic**
- Founder & CEO of **Pydantic Labs**

The Pydantic Stack

- **Pydantic Validation**: contracts for data
- **Pydantic AI**: production grade agent framework
- **Pydantic Logfire**: observability, evals, prompt management, LLMOps
- **Pydantic AI Gateway**: one endpoint, every model
- **Pydantic Monty**: Python Sandbox

**OpenTelemetry already exists.
Coding agents exist.**

So why did we write three SDKs and a CLI?

A great SDK gets out of the way

You should not have to learn it to use it

Two lines

```
from opentelemetry import trace
from opentelemetry.sdk.resources import Resource
from opentelemetry.sdk.trace import TracerProvider
from opentelemetry.sdk.trace.export import BatchSpanProcessor
from opentelemetry.exporter.otlp.proto.http import trace_exporter
from opentelemetry.instrumentation.fastapi import FastAPIInstrumentor
```

```
resource = Resource.create({'service.name': 'weather-app'})
provider = TracerProvider(resource=resource)
exporter = trace_exporter.OTLPSpanExporter(
    endpoint='https://logfire-us.pydantic.dev/v1/traces',
    headers={'Authorization': os.environ['LOGFIRE_TOKEN']},
)
provider.add_span_processor(BatchSpanProcessor(exporter))
trace.set_tracer_provider(provider)
FastAPIInstrumentor.instrument_app(app)
```

```
import logfire
```

```
logfire.configure()
logfire.instrument_fastapi(app)
```

- Same result
- The token comes from `.logfire/`, `LOGFIRE_TOKEN`, or an argument
- Repeat the left side twice more for metrics and logs; the right side already sends all three

Out of the way, in every language

```
import logfire

logfire.configure()
logfire.instrument_fastapi(app)
logfire.info('hello {name}', name=user)
```

```
import * as logfire from 'logfire'

logfire.configure()
logfire.info('hello {name}', { name })
```

```
let lf = logfire::configure().finish()?;
let _guard = lf.shutdown_guard();
logfire::info!("hello {name}", name = name)
```

- 34 `instrument_*` one-liners in Python: Pydantic AI, OpenAI, Anthropic, httpx, asyncpg, Django, Celery, Redis, MCP, Monty, ...
- Each one calls the OpenTelemetry instrumentation package for that library; instrumentation you already have keeps working
- JS wraps the OTel Node SDK and its auto-instrumentations; Rust is a `tracing` layer, so existing `tracing` spans are exported unchanged
- `send_to_logfire=False` and it exports to any OTLP backend

A great SDK is the agent's interface too

`configure()` and one `instrument_*` call per library is an API a model gets right.

The setup on the previous slide is not.

One binary, every ecosystem

```
➤ uv tool install logfire-cli # or: uvx logfire-cli
➤ npm install --global logfire-cli # or: npx logfire-cli
➤ pip install logfire # logfire 6 = logfire-sdk + logfire-cli
```

- One Rust binary. PyPI wheels for 17 targets, npm for 8. Neither Rust nor Python is needed to run it
- **For people:** `logfire signup`, `auth`, `init`, `project current`, `token read`
- **For scripts:** `--output json|jsonl`, `--no-input`, documented exit codes
- **For agents:** `logfire wizard` starts Claude Code or Codex in your repo with the Logfire setup skill; `logfire mcp ...` calls the hosted MCP server

Demo: auth + query

```
logfire auth
logfire mcp projects

logfire mcp query run \
  'SELECT span_name, duration, message FROM records ORDER BY start_timestamp DESC LIMIT 20' \
  --project demo

logfire --output json mcp query run --file slow.sql --project demo
```

- The query runs through the hosted MCP server, the same tool a coding agent calls
- Same binary and login for people, CI and agents; `--output json` for the last two

Prompts reach the agent through the same SDK


```
prompt = logfire.template_var(  
    name='prompt__weather_agent',  
    type=str,  
    default='Combine tool calls into a single code execution call. Answer in {{units}}.',  
    inputs_type=WeatherInputs, # a BaseModel with `units: str`  
)  
  
with prompt.get(WeatherInputs(units='celsius'), label='production') as resolved:  
    result = await weather_agent.run(question, instructions=resolved.value)
```

- A prompt is a managed variable named `prompt__<slug>`, with the same versions, targeting rules and rollouts as a feature flag
- No second vendor SDK: the import that sends traces also fetches prompts
- Other languages: Logfire serves **OFREP**, so any **OpenFeature** client can evaluate the same variables

A great SDK does what the protocol can't

OTLP carries finished spans. Your users are waiting on unfinished ones.

Pending spans



PLAIN OTLP: SPANS ARRIVE WHEN THEY END

- chat gpt-5.6 3.2s
- run_code 0.6s
- chat gpt-5.6 3.2s
- weather_agent run 7.6s
- GET /weather 8.0s

WITH PENDING SPANS: ROWS APPEAR WHEN THEY START

- GET /weather 8.0s
- weather_agent run 7.6s
- chat gpt-5.6 3.2s
- run_code 0.6s
- chat gpt-5.6 3.2s

- When a span starts, the SDK exports a zero-duration copy tagged `pending_span`; the backend replaces it when the real span arrives
- The copy is dropped if the real span is in the same export batch, so short spans are sent once

Values that look like your code

What an OTLP attribute can hold

```
"{\"id\": 1, \"email\": \"x@y.z\",  
  \"signed_up\": \"2026-01-01T09:30:00Z\"}"
```

A string, a number, a bool, or a flat list of those.

What Logfire shows

```
User(  
    id=1,  
    email='x@y.z',  
    signed_up=datetime(2026, 1, 1, 9, 30),  
)
```

The SDK sends a small schema next to the value, and the UI renders it as the Python you wrote.

- Works for Pydantic models, dataclasses, `datetime`, `Decimal`, `UUID`, `Enum`, `Path`, exceptions, DataFrames
- Other backends still receive the plain attribute

Demo: watch a trace arrive

```
logfire.configure(service_name='weather-app')
logfire.instrument_pydantic_ai()
logfire.instrument_monty()

app = FastAPI()
logfire.instrument_fastapi(app)

@app.get('/weather')
async def weather(q: str) -> dict[str, str]:
    result = await weather_agent.run(f"What's the weather in {q}, in Celsius?")
    return {'answer': result.output}
```

- One trace: HTTP request, Pydantic AI run, LLM call, Monty codemode execution, tool calls
- The rows are on screen before the model has answered

Thank you

Thanks to **Modal** for hosting



pydantic.dev/contact

@samuelcolvin · @pydantic